

K-Nearest Neighbors (KNN):

1. Model Overview

K-Nearest Neighbors (KNN) is a **non-parametric, instance-based (lazy learning)** algorithm used for both **classification** and **regression** tasks. It predicts outcomes based on the majority class (for classification) or average of nearest neighbors (for regression).

2. Key Aspects to Analyze for KNN

A. Assumptions

1. Similar data points exist close to each other in the feature space (locality assumption).
 2. The dataset is **representative and balanced**, with all relevant features captured.
 3. Distance metrics used (e.g., Euclidean, Manhattan) meaningfully represent similarity.
 4. Features are **scaled properly** (normalization or standardization is required).
 5. The data is **low-dimensional** — KNN suffers from the “curse of dimensionality.”
-

B. Limitations

1. **Computationally expensive** for large datasets — requires storing all data points.
 2. **Slow prediction time** — since distances are calculated for each query instance.
 3. **Sensitive to noisy or irrelevant features.**
 4. **Performance degrades** with high-dimensional data (many features).
 5. **Choice of ‘K’ and distance metric** significantly impacts accuracy.
 6. **Imbalanced classes** can bias the prediction toward the majority class.
-

C. Attributes / Input Features

1. Works best with **continuous and numerical** features.
 2. Categorical features can be handled using **appropriate encoding** and **distance functions** (e.g., Hamming distance).
 3. **Feature scaling** (e.g., Min-Max or Standard Scaler) is **mandatory**.
 4. Outliers should be minimized since KNN uses distance-based measures.
-

D. Internal Model Variations / Subtypes

1. **Weighted KNN** — assigns higher weights to closer neighbors.
 2. **Distance Metrics Variants:**
 - Euclidean Distance
 - Manhattan Distance
 - Minkowski Distance
 - Cosine Similarity
 3. **Approximate Nearest Neighbor (ANN)** — uses indexing structures for faster searches.
 4. **KNN Regression** — predicts continuous values using the mean (or weighted mean) of neighbors.
-

E. Hyperparameters

1. **K (Number of Neighbors):**
 - Low K → High variance (overfitting)
 - High K → High bias (underfitting)
 - Optimal K often found using **cross-validation** or **Elbow method**

2. **Distance Metric:** Defines how closeness is measured.
 3. **Weight Function:** Uniform vs distance-based weighting.
 4. **Algorithm:** 'auto', 'ball_tree', 'kd_tree', or 'brute' (in scikit-learn).
-

F. Performance Evaluation Metrics

For **Classification**:

- Accuracy, Precision, Recall, F1-score, ROC-AUC

For **Regression**:

- Mean Squared Error (MSE), Mean Absolute Error (MAE), R^2 Score
-

G. Use Cases

- Recommendation systems
 - Pattern recognition (e.g., handwriting, facial recognition)
 - Medical diagnosis (disease classification)
 - Credit risk and fraud detection
-

H. Optimization Tips

1. **Feature scaling** before model training.
2. **Dimensionality reduction (PCA)** for high-dimensional data.
3. **Remove outliers** and irrelevant features.
4. **Tune K** using validation curves or grid search.
5. Use **KDTree** or **BallTree** for large datasets to improve performance.

I. Model Interpretation

- **Transparent and explainable** — predictions are based on visible neighbors.
 - No internal coefficients or weights like linear models — relies purely on distance logic.
-

J. Summary Table

Aspect	Description
Type	Non-parametric, instance-based
Learning Type	Lazy (no training phase)
Key Parameter	K (number of neighbors)
Requires Scaling	Yes
Handles Outliers Well	No
Sensitive to Dimensionality	Yes
Ideal Dataset Size	Small to medium
Output Type	Classification or Regression