

# Linear Regression – Complete Notes


## 1. What is Linear Regression?

- Linear Regression is a **supervised machine learning algorithm** used for **regression tasks** (predicting continuous outcomes).
- The algorithm **learns from labeled data** (inputs XXX with known outputs YYY) to map inputs to output via a **linear function**.
- It assumes a **linear relationship** between input(s) and output: as input changes, output changes in a more-or-less proportional (linear) way.
- The form of that relationship is captured by a **straight line (in simple case)** or a hyperplane (in multiple dimensions).

### Example:

You observe that as students study more hours, their scores rise. You can use their study hours (input) to build a linear model to predict future exam scores (output).

- Independent variable (input) = hours studied
- Dependent variable (output) = exam score



## 2. Why Linear Regression is Important

Linear Regression has many appealing properties and uses:

- **Simplicity & interpretability:** It's easy to understand, visualize, and explain.
  - **Predictive power:** For tasks where the relationship is approximately linear, it gives good forecasts.
  - **Foundation for other models:** Many advanced methods build on the linear model concept.
  - **Efficiency:** Computations are cheaper (especially for small to medium dimensions).
  - **Widely used in practice:** It's a go-to method in economics, business, health, social sciences — wherever modeling relationships is needed.
  - **Insight into relationships:** The coefficients (slopes) tell how strongly each feature influences the target.
- 

## 3. Best-Fit Line in Linear Regression

To make predictions, linear regression fits a line (or hyperplane) that best “fits” the data points.

### Goal

- Find a line such that **predictions from the line are as close as possible** to the actual observed values.
- The closeness is measured via **errors (residuals)**, and the sum of errors (or sum of squared errors) is minimized.

### Equation of the Best-Fit Line (Simple Case)

For one independent variable  $x$ :

$$y = mx + b$$

$$Y = mx + c = (0.5) \cdot 5.5 + 32.76 =$$

$$y = m x + b$$




- $\hat{y}$  = predicted output

- $xxx$  = input
- $mmm$  = slope (gradient)
- $bbb$  = intercept (value of  $yyy$  when  $x=0$ )

In regression, we aim to choose  $mmm$  and  $bbb$  such that the line “fits” best.

### Interpretation

- **Slope ( $m$ ):** how much  $yyy$  changes (on average) for a unit increase in  $xxx$ .
- **Intercept ( $b$ ):** the predicted value of  $yyy$  when  $x=0$  (though sometimes  $x=0$  may not make sense in context).



## 4. Minimizing the Error: Least Squares Method

To decide which  $mmm$  and  $bbb$  are “best,” we need a criterion. The most common one is the **Least Squares criterion**:

- Compute **residual** for each data point:  

$$\text{residual}_i = y_i - \hat{y}_i$$
 (actual minus predicted)
- Form the **sum of squared residuals**:  

$$SSE = \sum (y_i - \hat{y}_i)^2$$
- We choose  $mmm$  and  $bbb$  to **minimize** that sum.

Why square residuals?

- To penalize large errors more heavily
- To make all errors non-negative (so negatives don't cancel positives)

Minimizing the SSE gives the **best-fit line** under assumptions of least squares.

---

## 5. Hypothesis Function

In machine learning terms, the model is called a **hypothesis**:

### 1. Simple Linear Regression

For one independent variable  $x$ :

$$h(x) = \beta_0 + \beta_1 * x$$

- $\beta_0 \rightarrow$  Intercept (bias term)
- $\beta_1 \rightarrow$  Coefficient (weight for the feature  $x$ )
- $h(x) \rightarrow$  Predicted value  $\hat{y}$

---

### 2. Multiple Linear Regression

For multiple features  $x_1, x_2, \dots, x_k$ :

$$h(x_1, x_2, \dots, x_k) = \beta_0 + \beta_1 * x_1 + \beta_2 * x_2 + \dots + \beta_k * x_k$$

- $\beta_0 \rightarrow$  Intercept (bias term)

- $\beta_1, \dots, \beta_k \rightarrow$  Coefficients (weights for each feature)
- $\hat{y} = h(x_1, x_2, \dots, x_k) \rightarrow$  Predicted value

This hypothesis function is what our learning algorithm tunes (finds best coefficients) to approximate  $y$ .

---

## 6. Assumptions of Linear Regression

For the model to be statistically valid and reliable, linear regression relies on several assumptions:

### 1. Linearity

- The relationship between inputs XXX and output YYY should be approximately linear.

### 2. Independence of Errors

- The residuals (errors) should not be correlated with each other.
- Particularly in time-series data, you must check for **autocorrelation**.

### 3. Constant Variance (Homoscedasticity)

- The residuals should have roughly constant spread (variance) across values of XXX.
- If the spread changes (fans out or narrows) — **heteroscedasticity** — it's problematic.

### 4. Normality of Errors

- The residuals should be approximately normally distributed (bell curve).
- Helps with inference (confidence intervals, hypothesis testing).

### 5. No Multicollinearity

- In multiple regression: the independent variables should not be too strongly correlated with each other.

- If two features are highly correlated, it's hard to separate their individual effects.

## 6. No Autocorrelation (for time series)

- Errors should not show patterns (especially in sequential/time data).

## 7. Additivity

- The effect of each feature on  $Y$  is additive (unless you explicitly model interactions).

If assumptions are violated, the estimates may be biased, inefficient, or misleading.

## 7. Types of Linear Regression

- **Simple Linear Regression** (Univariate): one predictor  $x$ .  

$$\hat{y} = \beta_0 + \beta_1 x$$
- **Multiple Linear Regression** (Multivariate): multiple predictors  $x_1, x_2, \dots, x_n$ .  

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Use simple when only one feature matters; use multiple when you have more explanatory features.

## 8. Cost / Loss Function for Linear Regression

The cost (or loss) function quantifies how far off the model's predictions are from the actual values.

- The common choice: **Mean Squared Error (MSE)**  

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

where  $n$  is number of data points.
- Some use a version with  $\frac{1}{2n}$  or  $\frac{1}{2}$  factors for convenience in derivative calculation.


Minimizing the cost gives us the best parameters.

## 9. Gradient Descent for Linear Regression

Gradient Descent is an optimization algorithm to find the parameters that minimize the cost function.

### How it works:

1. Initialize parameters ( $\beta_0, \beta_1, \dots, \beta_0, \beta_1, \dots$ ) to some values (often random).
  2. Compute predictions  $\hat{y}$  using current parameters.
  3. Compute the cost (MSE).
  4. Compute partial derivatives (gradients) of cost w.r.t each parameter.
  5. Update each parameter in the direction that reduces cost:
- $\alpha$  = learning rate (step size).
  - Repeat steps 2–5 until convergence (cost no longer decreases meaningfully).



Gradient descent helps find the global minimum (for convex cost) for linear regression.

---

## 10. Evaluation Metrics for Linear Regression

How do we measure how good the model is? Several metrics are used:

### 1. Mean Squared Error (MSE)

Measures the average squared difference between actual and predicted values:

$$\text{MSE} = (1 / n) \times \sum (y_i - \hat{y}_i)^2$$

---

### 2. Mean Absolute Error (MAE)

Measures the average absolute difference between actual and predicted values:

$$\text{MAE} = (1 / n) \times \sum |y_i - \hat{y}_i|$$

---

### 3. Root Mean Squared Error (RMSE)

The square root of MSE, interpretable in the same units as the target variable:

$$\text{RMSE} = \sqrt{\text{MSE}}$$

---

### 4. Coefficient of Determination ( $R^2$ )

Indicates how much variance in the target is explained by the model:

$$R^2 = 1 - (\text{RSS} / \text{TSS})$$

Where:

- $\text{RSS} = \sum (y_i - \hat{y}_i)^2 \rightarrow$  Residual Sum of Squares
  - $\text{TSS} = \sum (y_i - \bar{y})^2 \rightarrow$  Total Sum of Squares
  - $R^2$  ranges between 0 and 1; closer to 1 means better fit.
- 

### 5. Adjusted $R^2$

Adjusts  $R^2$  for the number of predictors to penalize adding irrelevant features:

$$\text{Adjusted } R^2 = 1 - [ (1 - R^2) \times (n - 1) / (n - k - 1) ]$$

Where:

- $n$  = number of data points
- $k$  = number of independent variables



- Helps prevent overestimation of model performance when adding unnecessary features.

These metrics help compare models and diagnose underfitting/overfitting.

---

## 11. 🧠 Regularization in Linear Regression

**Regularization is a technique used to prevent overfitting in regression models by adding a penalty term to the cost function.**

This penalty discourages the model from assigning excessively large weights to features.

When models overfit, they perform very well on training data but poorly on unseen (test) data.

Regularization introduces a bias–variance trade-off, improving generalization.

### 1. Lasso Regression (L1 Regularization)

Penalty Term:

$$\lambda \times \sum |\theta_j|$$

Lasso =  $\text{abs}(\text{ypred} - \text{ytrue})$

Modified Cost Function:

$$J(\theta) = (1 / 2m) \times \sum (\hat{y}_i - y_i)^2 + \lambda \times \sum |\theta_j|$$

Key Points:

- Encourages sparsity → Some coefficients become exactly zero.
- Performs feature selection by eliminating less important variables.
- Useful when you expect only a few features to have strong influence.

## 2. Ridge Regression (L2 Regularization)

Penalty Term:

$$\lambda \times \sum \theta_j^2$$

Modified Cost Function:

$$J(\theta) = (1 / 2m) \times \sum (\hat{y}_i - y_i)^2 + \lambda \times \sum \theta_j^2$$

Key Points:

- Penalizes large coefficients by shrinking them toward zero (but not exactly zero).
- Works well when all features contribute somewhat to the output.
- Reduces model variance but adds slight bias.

## 3. Elastic Net Regression (Combination of L1 and L2)

Penalty Term:

$$\alpha \lambda \times \sum |\theta_j| + (1/2)(1 - \alpha) \lambda \times \sum \theta_j^2$$

Modified Cost Function:

$$J(\theta) = (1 / 2m) \times \sum (\hat{y}_i - y_i)^2 + \alpha \lambda \times \sum |\theta_j| + (1/2)(1 - \alpha) \lambda \times \sum \theta_j^2$$

Key Points:

- Combines both Lasso and Ridge penalties.
- $\lambda$  (Lambda): Controls overall regularization strength.
- $\alpha$  (Alpha): Controls the mix between L1 and L2 regularization:
  - $\alpha = 1 \rightarrow$  Pure Lasso
  - $\alpha = 0 \rightarrow$  Pure Ridge
  - $0 < \alpha < 1 \rightarrow$  Mix of both

Elastic Net works well when there are multiple correlated features, where Lasso alone might arbitrarily drop one.

## 12. Python Implementation Example (Sketch)

Here's a high-level sketch (based on reference) of how linear regression can be implemented in Python from scratch:

1. **Import libraries**

`pandas, numpy, matplotlib`

2. **Load dataset**, split into training & testing sets.

3. **Define a class** `LinearRegression` with methods:

- `forward_propagation`: compute predictions  $\hat{y} = mx + c$
- `cost_function`: compute MSE
- `backward_propagation`: compute derivatives of cost wrt parameters
- `update_parameters`: apply gradient descent updates
- `train`: loop over many iterations, update parameters gradually, record loss

4. **Train** the model on the training data, plot the regression line and how it evolves.

5. **Use parameters** to predict on new/test data.

This implementation helps you deeply understand how linear regression mechanics work (beyond library calls).

---

## 13. Applications of Linear Regression

Linear regression is used across many fields. Some examples:

- **Finance & Economics**: Predicting stock prices, company performance metrics.
- **Real Estate**: Estimating property prices based on features (size, location, number of rooms).

- **Marketing / Business:** Forecasting sales given ad spend, seasonal factors, promotions.
- **Healthcare:** Predicting patient outcomes (e.g., length of hospital stay) from clinical variables.
- **Social Sciences / Psychology:** Modeling influence of factors on outcomes, e.g., income as a function of education, experience.

It helps both in predictive modeling and in **inferential analysis** (understanding which features matter and how).

---

## 14. Advantages & Disadvantages

### Advantages

- Simple to understand, implement & interpret
- Coefficients give meaningful insights
- Efficient to train, especially for moderate-size data
- Serves as a good baseline for more complex models
- Works well if the true relationship is close to linear

### Disadvantages / Limitations

- Assumes linearity — fails for non-linear relationships
  - Sensitive to outliers — a few extreme points can skew the line
  - Requires statistical assumptions (normality, independence, constant variance)
  - Multicollinearity among predictors causes instability
  - Can't capture complex patterns or interactions unless explicitly modeled
  - Overfitting / underfitting risk if model is too flexible or too simple
- 

## Summary & Flow

1. **Problem:** Identify target  $YYY$  (continuous) and input features  $XXX$ .
2. **Hypothesis:** assume model  $\hat{y} = \beta_0 + \sum \beta_i x_i$   
 $\hat{y} = \beta_0 + \sum \beta_i x_i$ .
3. **Cost function:** choose MSE to quantify error.
4. **Optimization:** use gradient descent (or direct methods) to find parameters that minimize cost.
5. **Check assumptions:** linearity, independence, homoscedasticity, normality, multicollinearity.
6. **Evaluate model:** use MSE, MAE, RMSE,  $R^2$ , adjusted  $R^2$ .
7. **Regularize** if overfitting threatens: ridge, lasso, elastic net.
8. **Deploy / Use model:** make predictions for new data; interpret coefficients to derive insights.